

A Guide for Bayesian Analysis in AD Model Builder

Cole C. Monnahan*,
Melissa L. Muradian, Peter T. Kuriyama

October 12, 2014

Abstract

The goal of this guide is to outline and describe the steps needed to conduct a Bayesian analysis in AD Model Builder. Included are general descriptions of Bayesian inference, priors, work flow and two built-in MCMC algorithms. We hope that the guide will enable users to take full advantage of the features of ADMB that are built-in, but not well-documented.

Contents

1	Introduction	2
2	Bayesian inference	2
3	Priors	3
4	Markov chain Monte Carlo (MCMC) in ADMB	3
4.1	Work flow	3
4.2	MCMC Phases	4
4.3	Output files	5
4.3.1	Meta data: The hst file	5
4.3.2	Parameter draws: The psv file	5
4.3.3	Derived quantity draws	5
4.4	Convergence diagnostics	6
4.5	Resuming a chain	6
4.6	Starting values	8
5	Metropolis-Hastings MCMC	8
5.1	Algorithm	8
5.2	MCMC Arguments	8
5.3	Console output	9
5.4	Example MCMC	9
5.5	Tuning the MH algorithm	10
5.5.1	Thinning rate	10
5.5.2	Optimize the acceptance rate	11
5.5.3	mcpb	11
5.5.4	mcrb	12
5.5.5	User-specified covariance matrix	14

*corresponding author; monnahc@uw.edu

6	Hybrid MCMC	14
6.1	Algorithm	15
6.2	Arguments	16
6.3	Tuning the hybrid algorithm	17
7	An example with non-linear posterior	18

1 Introduction

AD Model Builder is a fast, powerful software for fitting complex, nonlinear statistical models[1]¹. The “AD” refers to automatic differentiation, a powerful method for calculating derivatives of the likelihood to aid in quickly fitting models. While ADMB is designed for maximum likelihood inference, it also contains the capability of Bayesian inference by numerically sampling from the posterior distribution of a model. This option may be desired, however the capabilities and options available may be unclear, particularly for new users. This guide is an attempt to clarify how to run Bayesian analyses for models built with ADMB.

ADMB models, which are compiled C++ executables with a wide variety of runtime arguments, can be developed in your favorite text editor, or the packaged ADMB IDE. Compiling and running executables can be done through the terminal, or R users can take advantage of the `R2admb` package available on CRAN. This package creates an R to ADMB interface useful for compiling and running models, and reading results back for plotting or post-processing. This guide assumes the user is comfortable building, compiling and running ADMB models to get maximum likelihood estimates of parameters and standard errors. See the official documentation on the ADMB website for more information if you are unfamiliar with this aspect of the software. We further assume a base level of knowledge of the Bayesian philosophy and numerical techniques, although we do provide a brief introduction and description of the algorithms available to the user.

This guide began because we use ADMB to run Bayesian models for our own research, and we needed to look into the source code to fully understand some of the behavior and options available². The writeup arose as a way to cement our understanding of the internal workings of ADMB, and as a way to give back to a free product. This guide is unofficial, informal, and most importantly open source; we welcome and encourage users to contribute to this guide to further document the Bayesian capabilities of ADMB, or other parts of the ADMB project. The development contents of this document available on a Github repository maintained by Cole Monnahan. Improvements can be merged directly to the Rnw file (`knitr`), and suggestions submitted as issues on the Github page.

2 Bayesian inference

A full description of Bayesian inference is beyond the scope of this guide, and we encourage users unfamiliar with it to seek more appropriate sources. In general, Bayesian inference is a method of statistical inference used to quantify the full range of uncertainty around both models and parameter values, and formally incorporate prior information with observed data. It uses the same likelihood function as the maximum likelihood method of inference, but also explicitly combines prior information, typically as posterior probability distributions from previous studies, to create a probability distribution for the parameters (called the “posterior”).

Markov chain Monte Carlo (MCMC) is a common numerical algorithm used to sample from posteriors of complex and high-dimensional probability distributions, particularly where an analytical solution is intractable. ADMB implements two different MCMC algorithms: the ubiquitous Metropolis-Hastings and a Hamiltonian (or “hybrid”) sampler, both briefly described in this guide.

The goal of MCMC algorithms is to draw samples from the posterior distribution, from which inference (e.g. parameter estimates and credible intervals) can be done. In a sense it can be thought of as a numerical integration technique. While MCMC is very flexible, it can be highly inefficient, needing thousands or millions of function calls to get a sufficiently well-defined posterior for reliable inference. As such, the idea

¹see <http://admb-project.org/>

²See the API for a convenient way to explore the source code

of tuning the MCMC chain to be more efficient can be important to reduce runtime, and we discuss this in the context of ADMB for both MCMC algorithms. If the chain is not tuned properly, samples taken from it could lead to invalid inference. **The responsibility is on the user to verify the validity of samples from a MCMC chain.** The R package CODA is a commonly used tool for checking MCMC diagnostics such as autocorrelation. Think of ADMB as a tool to generate posterior samples, and these must then be validated and inference calculated by the user.

3 Priors

As with any Bayesian analysis, priors are a key component of the posterior for ADMB and should be explicitly declared in the `.tpl` file. The posterior density is proportional to the product of the likelihood and priors. ADMB works in the negative log space, so the contribution of priors is incorporated into the model by adding the negative log of the prior density to the objective function. If priors are added, the minimum point is no longer the maximum likelihood estimate, but rather what we informally refer to as the posterior mode – i.e. the parameters corresponding to the highest point of the posterior (or lowest point on the $-\log$ posterior surface). The difference between the two is subtle, but there are important philosophical differences: with Bayesian analysis ADMB is integrating across the posterior parameter space, while the likelihood framework uses minimization of the likelihood to get parameter estimates and asymptotic assumptions for confidence intervals.

When running an MCMC in ADMB, users are **not** forced to explicitly define and incorporate priors into their model. In this case “implicit” priors are still affecting the posterior. In particular, unbounded parameters without explicit priors have infinitely-broad uniform priors (an example of an “improper” prior because it will not integrate to 1). Another example is a variance parameter estimated in log space and then exponentiated within the model to keep it above zero. An implicit uniform prior on the log parameter will provide different prior information than if the parameter were on the natural scale. The implications of these issues are contentious and certainly beyond the scope of this guide. We caution users to carefully consider and explicitly acknowledge their priors – ADMB will not force you to.

Bounding a parameter automatically incorporates a uniform prior without changing the objective function calculation, since values outside the bounds have zero probability. Other forms of prior distributions must be added manually with, or without, any constants. For instance, a normal ($N(\mu, \sigma^2)$) prior on a parameter a could be added as `f+=pow((a-μ),2)/(2*pow(σ,2));`, where the constants have been dropped and `f` is the objective function.

4 Markov chain Monte Carlo (MCMC) in ADMB

MCMC is the only form of built-in Bayesian analysis available to users of ADMB, although a wide variety of algorithms exist for other software platforms. John Kruschke has a simple example that illustrates the mechanics of the Metropolis-Hastings algorithm[2]. A politician lives on a long chain of islands. The politician’s goal is to stay in the public eye by traveling to each island proportional to their relative population. In other words, he will spend more time on highly populated islands and less time on less populated islands. He flips a fair coin to propose a move to an island east or west of his current location. If the proposed island has a higher population than his current island he moves. If the proposed island has a smaller population than his current island he will only move with a probability P_{prop}/P_{cur} where P_{prop} is the population of the proposed island and P_{cur} is the population of the current island.

MCMC algorithms work in this same general way. The algorithm starts at a set of parameters, and then proposes a new parameter set. The algorithm moves to the proposed state, or not, depending on its density relative to the current state. Just like the politician, the algorithm will spend most of the time in states with high densities.

4.1 Work flow

The R package `R2admb` contains some useful functions for streamlining the work flow of MCMC in ADMB by using R. For users who wish to have more manual control, the following steps outline a typical work flow.

1. Build, run, and verify your ADMB model. This model should explicitly include the contribution of the priors to the objective function, such that the ADMB point estimates are the posterior mode, rather than a maximum likelihood estimate.
2. Run an MCMC using the command line argument `-mcmc N -mcsave Nsave` (see Table 2 for more runtime options). The unsaved draws are discarded, leaving a total of $N_{\text{out}} = N/N_{\text{save}}$ saved draws. For example, `-mcmc 1e6 -mcsave 2000` will run 1 million draws but only save every 2000th (i.e. “thin” the others), for a total kept of $N_{\text{out}} = 500$.
3. After completion, run the model again with argument `-mceval`. This command tells ADMB to loop through the saved iterations (in the `.psv` file) and execute in the `mceval_phase()`.
4. Pull results into R, Excel or other program to ensure the sample is sufficiently thinned, either visually or with tools using, for example, the `CODA` package.
5. If necessary, rerun the chain with more thinning, drop the first part of the chain as a “burn-in,” or run longer for more saved draws.
6. Make whatever Bayesian inference is desired using the N_{out} independent samples. For instance, calculate the median and 95% credible interval for parameters or derived quantities (functions of parameters), or plot pairwise parameters to explore correlation.

4.2 MCMC Phases

ADMB is designed with two phases that are used to produce MCMC output: (1) the `mcmc` phase and (2) `mceval` phase. In our experience, the use of these phases is not common among other MCMC software and may be a source of confusion for new ADMB users. However, they provide a powerful and efficient framework for MCMC analyses.

Most people are already familiar with the operations performed in the `mcmc` phase. This is where ADMB uses an algorithm to draw samples (parameter vectors) from the posterior. New states are proposed using one of two built-in algorithms: see section 5 and 6. During this phase, the N_{out} saved parameter values are written to a `.psv` file (described below). Note that if `-mcsave Nsave` is not specified at runtime, ADMB will run the MCMC **but no values will be saved**.

The `mceval` is an optional phase that is designed to be run after the `.psv` file has been produced. During this phase, ADMB loops through the N_{out} parameter combinations in the `.psv` file and reruns the `PROCEDURE_SECTION`. This phase is extremely powerful because it allows the user to minimize runtime by partitioning calculations into two groups: those that affect the objective function (i.e. posterior calculations) and those that do not. Calculations done for discarded draws (which typically comprises most iterations) simply slow down the analysis, and as such should be minimized as much as possible. For example, a user may want to extrapolate (e.g. project a value into the future, given current parameters) or calculate values derived from the parameters and intermediate values. Placing these calculations inside an `mceval_phase` clause means they will only be performed on each saved draw. Thus, an analysis can be made to run faster by minimizing total calculations in the `mcmc` phase. In practice, some chains need to be thinned significantly more than 1 in 1000, so the runtime saved can be substantial, especially if the calculations in the `mceval_phase` are time consuming.

While the `mceval` phase was designed specifically for MCMC analyses, it can be co-opted for use in other types of analyses. In essence it is a convenient framework in which to get ADMB to quickly evaluate arbitrary sets of parameters, while only initializing once. Examples of alternative uses are the SIR algorithm, evaluating a grid of points for plotting and exploration of the posterior surface, or trying random parameter sets to investigate local minima. Getting ADMB to evaluate these parameter sets is as simple as writing them to the `.psv` file and then executing ADMB with the option `-mceval`. See section (4.3) for details on how to do this.

4.3 Output files

4.3.1 Meta data: The hst file

A file named `<model name>.hst` is produced which contains the marginal distribution of each object of type `sdreport` displayed as two columns under the header of the associated `sdreport` variable. This distribution is created by binning the observed values of each `sdreport` variable, over the entire MCMC chain, into a histogram and then scaling the area under the curve to equal 1. In the `.hst` file, the first column reports the mid-point of each bin and the second column reports the associated density.

The structure of the `.hst` file is organized so that if a chain is restarted using the `-mcr` option (see section 4.5), all values needed to bin the subsequent MCMC samples into the original histogram are automatically read, so that the values in the final `.hst` file reflect the observed values, and therefore the marginal distribution, of the complete chain. The first element reported in the `.hst` file is the total number of iterations of the MCMC chain, N_{mcmc} , used to create the reported marginal distribution, and the second element is the final step-size scaling factor used for the jump function. The remainder of the elements are only useful for constructing the reported histogram.

An important note about using the `.hst` file: The reported means and standard deviations are necessary to the internal algorithm that creates each `sdreport` variable's histogram and are calculated using only a small number of the initial iterations of the chain. Therefore these values are inappropriate to report as model results. The user has two ways to obtain the posterior median and credible intervals for each `sdreport` variable. The user can manually export the marginal distribution from the `.hst` file into an input file external to ADMB (e.g. `.csv` to R) and calculate the quantiles, etc. Alternatively, the user can write each `sdreport` variable's (saved or total) draws during the `mceval` phase to a `.csv` by declaring an IO object. More information can be found in the derived quantity draws section 4.3.3. This `.csv` can then be read into R and manipulated to produce a histogram of the observed values.

If the user is writing the saved `sdreport` variables to file, then the only pieces of information in the `.hst` file that are useful are the first two elements, N_{mcmc} and the final step-size scaling factor.

4.3.2 Parameter draws: The psv file

During the `mcmc` phase, saved parameter values, in bounded space, are written to a binary file called `<model name>.psv`. This file can be read into R using the following commands:

```
psv <- file("<model name>.psv", "rb")
nparams <- readBin(psv, "integer", n = 1)
mcmc <- matrix(readBin(psv, "numeric", n = nparams * Nout), ncol = nparams,
               byrow = TRUE)
close(psv)
```

The first element in the `.psv` file is the number of active parameters in the model, which then tells R how to parse the following elements into parameter values. Note that the value of N_{out} in `nparams*Nout` depends on N_{mcmc} and `mcsave` and must be specified manually. This is the main file that was designed to be used to extract MCMC draws from ADMB. However, this file only contains parameter values and not derived quantities or other quantities of interest (e.g. *MSY* or biomass trajectories) which often are of interest.

4.3.3 Derived quantity draws

Often the posterior distribution for quantities other than the parameters are desired. Examples of derived quantities are: functions of parameters, properties of the model, or model projections/extrapolations.

A simple way of extracting this information is to bypass the `psv` file altogether and use a C++ function to write a `.csv` file containing whichever quantities are desired. This can be accomplished inside the ADMB `.tpl` file with just a few lines of code. Inside the `DATA_SECTION` section use the following code to create an IO object that writes values to a `.csv` file, similar to the function `cout` which prints to screen.

```
!!CLASS ofstream MCMCreport("MCMCreport.csv", ios::trunc);
```

Then, inside the `PROCEDURE_SECTION` the function can be used to write both parameters, derived quantities (`sdreport` variables, etc.), or other information about the model. For example, during the `mceval` phase we may want to write the model parameters, model objective function, and function of parameters to file. This can be accomplished as follows:

```
if(mceval_phase()){
  if(header==1) {
    MCMCreport << "a,b,f,ab" << endl;
    header=0;
  }
  MCMCreport << a <<"," << b << "," << f << "," << ab << endl;
}
```

In this case the parameters a and b are saved, along with the negative log-likelihood (f) and product of the model parameters (ab). The `MCMCreport` object is used just like `cout` and is executed only during the `mceval` phase so that only saved values are written to the file as `ADMB` iterates through each saved parameter during that phase. The variable `header` is set to 1 in the `GLOBALS_SECTION` as `int header = 1;`. Naturally, this strategy can be used anywhere in the procedure section, and this may be a useful diagnostic tool in some situations. The `csv` is recreated at each execution if `ios::trunc` is used. Alternatively, new draws are appended to the bottom of the `csv` with `ios::app`; this option is particularly useful when restarting an MCMC chain. If the append option is used, but the corresponding `csv` doesn't exist in the working directory, the file is created and filled.

4.4 Convergence diagnostics

There exist important convergence diagnostics to assess whether an MCMC chain has not “converged” (roughly speaking this means it is producing independent samples from the posterior). There are a few key diagnostics to check: 1) whether enough of the initial samples have been discarded as “burn-in”, 2) how much auto-correlation exists in the saved samples. Using the saved parameter values, histograms and trace plots for each parameter should be created and reviewed, perhaps using the `CODA` package [3]. Invalid samples (e.g. not converged, stuck on bounds, during tuning phase) can lead to invalid inference, so we stress the importance of diligently checking properties of the chain outside of `ADMB`.

The first few iterations of a chain should be discarded as “burn-in”; the period where the algorithm is adjusting the jump function in order to best explore the posterior space (figure). A general rule of thumb of a 20%-50% burn-in is often recommended. However, this is typically not required for `ADMB` since it starts from the posterior mode (a region of high probability density, see 4.6), whereas many other MCMC software platforms start from arbitrary places and the chain spends a long time moving to the area of high density. At a minimum values during which the chain is being tuned should be discarded. After the samples judged to be from the burn-in period have been manually discarded, trace plots should be trendless, in that consecutive samples should not be correlated.

If autocorrelation exist in the histogram even after the burn-in has been removed, then the chain should be run over again with a greater thinning rate, e.g. from `-mcsave 10000` to `-mcsave 50000`. This step is crucial to obtain a final chain of independent samples. If there still exist strong trends in the trace plots, there is correlation between two or more parameters. Correlation between parameters can be seen using the `pairs()` plotting function on an `mcmc` object in R.

More information on these concepts and examples are given in sections 5.5 and 6.3.

4.5 Resuming a chain

The `MC resume` command, `-mcr`, allows the user continue sampling along a previously run MCMC chain. Say that the user has run a short ($N = 10,000$) chain, checked autocorrelation diagnostics, and concluded that the thinning rate is appropriate. Rather than restarting a new chain and discarding these 10,000 iterations, the user can pick up an existing chain where it left off with the command `-mcr`. Alternatively, if further thinning is needed, the user can continue an existing chain and thin outside of `ADMB` after the chain is complete.

Resuming a chain is simple, and only requires adding `-mcr` to an MCMC statement. Let's save ten values from a chain using the "simple" model (a contrived linear model packaged with ADMB), saving every 20th value with a seed of 30.

```
simple -mcmc 200 -mcsave 20 -mcseed 30
```

Now say that we want to resume the chain, still saving every 20 values until we have 20 saved parameter draws. In order to save ten additional draws we will run 200 more additional iterations by adding `-mcr` (see Section 2.1 in the ADMB manual).

```
simple -mcmc 200 -mcsave 20 -mcr
```

One nuance of `-mcr` is that starting then resuming a chain does not result in the same samples as a long chain. For example if we run 400 iterations of a chain, saving every 20, and setting the seed at 30 with the command:

```
simple -mcmc 400 -mcsave 20 -mcseed 30
```

we end up with different values from the started and resumed chain (Table 1). Users should note the differences between the last half of the "resumed" chains and the "long" chains. This disparity can have important implications for reproducibility.

Resuming chains with `-mcr` can reduce computation time. Adding the command `-noest` can save further computation time as minimization is unnecessary for resumed chains. Adding the command `-nosdmcmc` can also save computation time, but the original chain must have been run with `-nosdmcmc`. Otherwise, ADMB will try to read non-existent histogram data (found in the `.hst` file).

a-first	b-first	a-resumed	b-resumed	a-long	b-long
1.91	4.08	1.91	4.08	1.91	4.08
2.00	4.78	2.00	4.78	2.00	4.78
2.08	3.40	2.08	3.40	2.08	3.40
1.85	4.35	1.85	4.35	1.85	4.35
1.88	4.31	1.88	4.31	1.88	4.31
1.86	4.94	1.86	4.94	1.86	4.94
1.89	2.18	1.89	2.18	1.89	2.18
2.62	0.78	2.62	0.78	2.62	0.78
2.16	2.99	2.16	2.99	2.16	2.99
1.94	3.49	1.94	3.49	1.94	3.49
0.00	0.00	0.33	-1.25	2.10	2.71
0.00	0.00	2.31	1.02	1.90	4.38
0.00	0.00	1.93	3.36	1.91	4.00
0.00	0.00	1.73	5.22	2.08	4.00
0.00	0.00	1.59	6.24	1.59	5.48
0.00	0.00	2.16	3.08	1.72	5.47
0.00	0.00	1.83	3.24	1.66	5.69
0.00	0.00	1.95	4.12	2.26	3.02
0.00	0.00	1.76	4.55	1.74	4.69
0.00	0.00	1.82	5.21	2.38	2.51

Table 1: Parameter draws from the "simple" model. The "first" columns are saved parameter draws from a short chain, $N_{\text{out}} = 10$. The values from the "resumed" chain were created using the `-mcr` option, where they pick up from where the "first" chain left off, $N_{\text{out}} = 20$. The values shown from the "long" chain were completed in a single run (not resumed with `-mcr` and set to match the total number of iterations and thinning rate from the "reduced" chain).

4.6 Starting values

For many Bayesian software platforms, the MCMC algorithms are started at user-specified or arbitrary places. ADMB has the advantage that it can robustly estimate the posterior mode and the covariance at that point. This information is very valuable in initializing the MCMC chain.

Specifically, an MCMC chain starts from the posterior mode and uses the estimated covariance matrix in its proposed jumps (see the algorithm sections below). As such, ADMB chains typically do not need a long period to reach areas of high density.

The user can specify starting values via the command line option `mcpin <file>` which must point to a file with starting values for parameters.

5 Metropolis-Hastings MCMC

The default MCMC algorithm used by ADMB is the Metropolis-Hastings (MH) algorithm³. This algorithm has been around for decades, is simple to implement, and used widely.

This algorithm will be most efficient when the posterior surface mimics a multivariate Normal distribution.

5.1 Algorithm

Let

f = the ADMB objective function
 c = an unknown normalization constant
 X_{cur} = current parameter vector
 X_{prop} = a proposed parameter vector
 U = a randomly drawn uniform value in $[0,1]$

Then

$$X_{\text{new}} = \begin{cases} X_{\text{prop}} & \text{if } U \leq \frac{cf(X_{\text{prop}})}{cf(X_{\text{cur}})} \\ X_{\text{cur}} & \text{otherwise} \end{cases} \quad (1)$$

The proposal (or “jump”) function proposes new parameter vectors given the current set. The default behavior for ADMB is to use a multivariate normal distribution⁴ centered at the current vector:

$$X_{\text{prop}} \sim MVN(X_{\text{cur}}, \Sigma)$$

where Σ is the covariance matrix obtained by inverting the Hessian at the posterior mode.

5.2 MCMC Arguments

ADMB has a suite of arguments available to the user at run time. We cover the most commonly used ones in table 2, and refer the user to the ADMB manual and other documentation for further options.

³Technically it has a symmetric proposal, a special case known as the Metropolis algorithm.

⁴Technically a bounded multivariate normal, operating in a transformed parameter space via the Cholesky decomposition of the covariance matrix, but we ignore that here for simplicity and clarity.

<code>-mcmc N</code>	Run N MCMC iterations
<code>-mcsave N</code>	Save every N th MCMC iteration
<code>-mcscale N</code>	Rescale step size for first N iterations
<code>-mcmult N</code>	Rescale the covariance matrix
<code>-mcpin <file></code>	Start algorithm from values in <code> file </code>
<code>-mcrb N</code>	Reduce high parameter correlations (see 5.5.4)
<code>-mcprobe X</code>	Use a fat-tailed proposal distribution (see 5.5.3)
<code>-mcdiag</code>	Use a diagonal covariance matrix
<code>-mcnoscale</code>	Do not scale the algorithm during
<code>-mcscale N</code>	Scale the algorithm during the first N iterations
<code>-mcu</code>	Use a uniform distribution as proposal distribution.

Table 2: ADMB runtime arguments for the Metropolis-Hastings MCMC

5.3 Console output

When running an MCMC chain from the console, the output similar to the following will appear.

```
Initial seed value -36519
-14.9642 -14.9642
  mcmc sim 1  acceptance rate 0 0
-15.4843 -15.4843
  mcmc sim 201 acceptance rate 0.517413 0.52
increasing step size -1.32761
-15.7798 -15.7798
  mcmc sim 401 acceptance rate 0.428928 0.34
```

ADMB prints the random seed used, and then the negative objective function value twice. In this case scaling is increased after the first 200 iterations (proposing more distant values). This is followed by the log of the determinant of the Cholesky decomposed covariance matrix. In some cases this information may be useful for troubleshooting, but for most practical purposes it can be ignored.

5.4 Example MCMC

We demonstrate these concepts using the “simple” model packaged with ADMB, which is a contrived two parameter linear model ($y = ax + b$). First we run the model without any thinning by specifying `mcsave=1` (figure 1). The MCMC chain is clearly autocorrelated from the acf plots and the traces (not show). This chain needs to be tuned to achieve independent samples.

5.5 Tuning the MH algorithm

The goal of “tuning” the algorithm is to obtain independent samples from the posterior, with as few of calculations as possible (to keep runtime manageable). For the MH algorithm, there are several ways the user can tune a chain.

5.5.1 Thinning rate

For the MH algorithm, the most important tuning option available to the user is the saving rate (the inverse of the thinning rate). This is the rate at which parameters are saved, such that thinning is effectively discarding draws. This tuning option is critical since this algorithm generates autocorrelated parameters by design.

The user controls the thinning rate by the argument `mcsave N`. If $N = 1$, as in figure 1, every single draw is saved (none are thinned out). As we saw with that example, the autocorrelation is high, suggesting the need to thin more (save fewer).

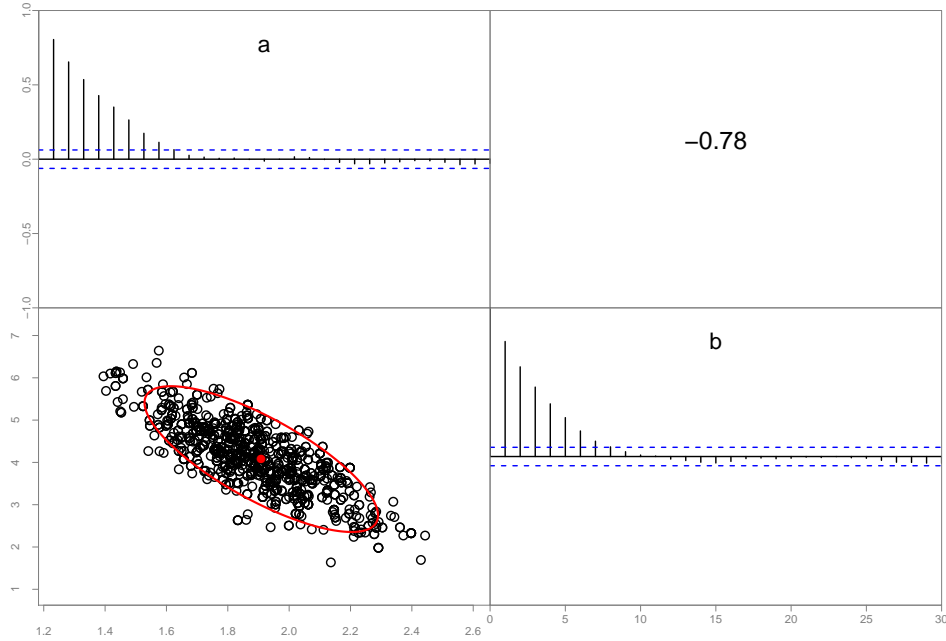


Figure 1: The samples from a simple model with `mcsave=1`. Note the high autocorrelation of both parameters. The red ellipses show the estimated pairwise parameter covariances, and the red point the MPD.

We now rerun the chain with `mcsave 100` (figure 2), by increasing the total samples by 100 and saving every 100th. This helps reduce the autocorrelation and produces independent draws from the posterior of interest.

5.5.2 Optimize the acceptance rate

Studies have shown that there is an optimal range for acceptance rate for the MH algorithm (e.g. [4]). If the proposal distribution generates values too close to the current state, the chain will accept them (high acceptance rate) but explore the posterior slowly and need more thinning. Alternatively, if proposals are too far away into regions of low density (low acceptance rate) the chain will not explore the space. The optimal acceptance rate varies by model, but is roughly 40%. The general advice is to tune the proposal distribution to achieve an efficient acceptance rate.

ADMB accomplishes this by “scaling” the covariance matrix up or down, depending on the current acceptance rate, during the first part of the chain. Scaling the covariance matrix down produces proposed sets closer to the current set, and vice versa for scaling up. By default, it scales during the first 500 iterations (and prints this to screen), but the user can specify this with `mcscale N` or turn off scaling with `mcnoscale`. ADMB rescales the covariance matrix every 200 iterations until the acceptance rate is between 0.15 and 0.4, or the scaling period is exceeded.

In practice, the defaults work for many models, but the user may want to extend the scaling period for some models. Draws from this tuning phase should be discarded as part of the burn-in.

5.5.3 `mcprobe`

For some models, there may be concern of being “stuck” in a local minimum and simply never proposing a value far enough away to escape it and find other regions of high density. Obviously this problem would present issues for maximum likelihood inference as well. ADMB has a built-in algorithm which modifies the default proposal distribution so it occasionally proposes very distant parameters (i.e. “probes”)⁵. The

⁵Previous versions of ADMB called this `mcgrope` – this is now deprecated

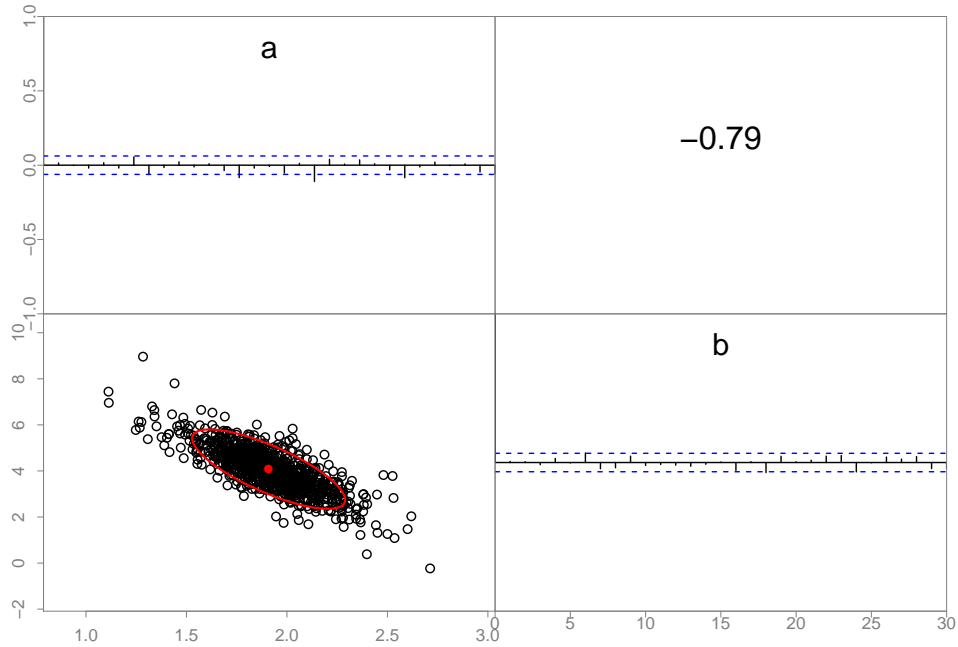


Figure 2: The samples from a simple model with `mcsave=100`. Note the negligible autocorrelation of both parameters. The red ellipses show the estimated pairwise parameter covariances, and the red point the MPD.

`mcp` **probe** **X** argument initiates this option.

The modified proposal distribution is a mixture distribution of normal and Cauchy distributions. The argument **X** controls how the two distributions are mixed, with larger values being more Cauchy (fatter tails, larger jumps). The range of valid inputs is 0.00001 to 0.499, and if no value is supplied a default of 0.05 is used⁶

Figure 3 shows the shape of the proposal distribution. Note the extremely fat tails, and that the standard proposal itself is bounded.

5.5.4 `mcrb`

The `-mcrb N` option (which stands for rescaled bounded) alters the covariance matrix used to propose new parameter sets in the MH algorithm. Its intended use is to create a more efficient MCMC sampler so the analyses run faster.

The option will be most effective under circumstances where the correlation between parameters at the posterior mode is higher than other regions of the parameter space. In this case, the algorithm may make efficient proposals near the posterior mode, but inefficient proposals in other parts of the parameter space. By reducing the correlation using `mcrb` the proposal function may be more efficient on average across the entire parameter space and require less thinning (and hence run faster).

⁶See the function for more on how the distribution is created




Figure 3: Example of fat-tailed proposal values for a parameter for the `mcprobe` option compared to the default proposal.

The `mcrb` option is a set of calculations performed on the original correlation matrix, as follows.

$$\begin{aligned}
 \Sigma_{\text{old}} &= \begin{bmatrix} 1 & \cdots & \rho_{1,n} \\ \vdots & \ddots & \vdots \\ \rho_{n,1} & \cdots & 1 \end{bmatrix} && \text{The original correlation matrix} \\
 \mathbf{L} &= \begin{bmatrix} 1 & \cdots & 0 \\ \vdots & \ddots & \vdots \\ L_{n,1} & \cdots & L_{n,n} \end{bmatrix} && \text{Lower Choleski decomposition of } \Sigma_{\text{old}} \\
 \hat{\mathbf{L}} &= \begin{bmatrix} 1 & \cdots & 0 \\ \vdots & \ddots & \vdots \\ L_{n,1}^{N/10} & \cdots & L_{n,n}^{N/10} \end{bmatrix} && \text{Raise elements to power user supplied } N \\
 \tilde{\mathbf{L}} &= \begin{bmatrix} 1 & \cdots & 0 \\ \vdots & \ddots & \vdots \\ \frac{\hat{L}_{n,1}}{|\hat{L}_{n,\cdot}|} & \cdots & \frac{\hat{L}_{n,n}}{|\hat{L}_{n,\cdot}|} \end{bmatrix} && \text{Normalize rows of } \hat{\mathbf{L}} \\
 \Sigma_{\text{rb}} &= \tilde{\mathbf{L}}\tilde{\mathbf{L}}^T && \text{Calculate new correlation matrix}
 \end{aligned}$$

By working with the Cholesky decomposition of the correlation matrix, the algorithm ensures that the rescaled bounded matrix used in the MCMC remains a valid correlation matrix. The impacts on a correlation matrix can be difficult to anticipate, but fortunately ADMB writes the resulting covariance matrix to the file `corrtest` (a text file with no ending) under the label `modified S`. The user can plot this against the estimated covariance to visually gauge the impact of the `mcrb` algorithm.

If poor performance is suspected to be caused by correlations that are too high, the `mcrb` option provides a quick, convenient way to try a reduced correlation matrix in the algorithm.

5.5.5 User-specified covariance matrix

The `mcrb` option is a quick way to try lowering the correlation between some parameters. A more flexible (and transparent) option is to provide ADMB with a covariance matrix specified by the user. By examining preliminary chains or using information gained from other simple models, it may be possible for the user to define a covariance matrix that is more efficient than the estimated one, or modified via the built-in options. We demonstrate the improved performance using this option in the next section.

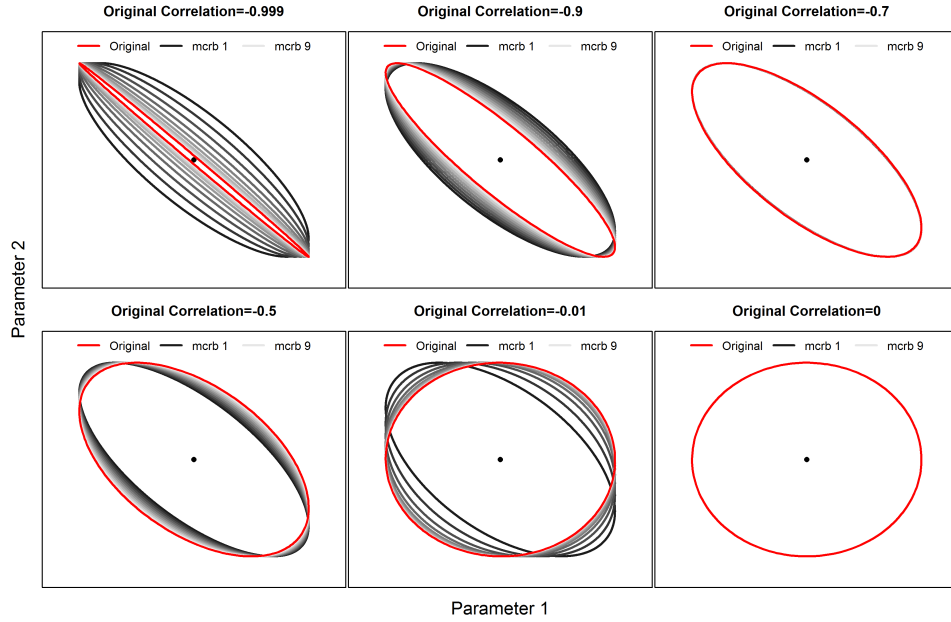


Figure 4: The effect of mcrb on a variety of correlations between two hypothetical parameters. Note that the effect of setting $N = 9$ depends on the original correlation.

Providing user-supplied covariance is a more advanced technique that is not built in to ADMB. It requires a slightly more detailed understand of the underlying output file structure of ADMB. The steps to do this can be found in a separate document on the ADMB website. The key is that the `admodel.cov` file is in the unbounded space and thus the user specified covariance matrix must be converted using the following function or something similar:

```
write.admb.cov <- function(cov.user, model.path=getwd()){
  temp <- file.exists(paste0(model.path, "/admodel.cov"))
  if(!temp) stop(paste0("Couldn't find file ",model.path, "/admodel.cov"))
  temp <- file.copy(from=paste0(model.path, "/admodel.cov"),
                    to=paste0(model.path, "/admodel_original.cov"))

  wd.old <- getwd()
  on.exit(setwd(wd.old))
  setwd(model.path)
  ## Read in the output files
  results <- get.admb.cov()
  scale <- results$scale
  num.pars <- results$num.pars
  if(NROW(cov.user) != num.pars)
    stop(paste0("Invalid size of covariance matrix, should be: ", num.pars,
               "instead of ",NROW(cov.user)))
  cov.unbounded <- cov.user/(scale %o% scale)
  ## Write it back to file
  file.new <- file(paste0(model.path, "/admodel.cov"),"wb")
  on.exit(close(file.new))
  writeBin(as.integer(num.pars), con=file.new)
  writeBin(as.vector(as.numeric(cov.unbounded)), con=file.new)
  writeBin(as.integer(results$hybrid_bounded_flag), con=file.new)
  writeBin(as.vector(scale), con=file.new)
}
```

We note that the user is responsible for checking whether the covariance matrix passed is valid. In R this can be done with the `matrixcalc` package:

```
matrixcalc::is.positive.definite(x, tol=1e-8)
```

If the covariance matrix is not valid, ADMB will throw an error and perhaps proceed but the output from the MCMC would likely be unusable.

6 Hybrid MCMC

The “hybrid” option in ADMB is an implementation of an MCMC algorithm based on Hamiltonian dynamics. Here we provide a simplified overview⁷ of the algorithm, with the aim of providing users an intuition about its behavior and properties and how to use it within ADMB. This section is based on [5], which provides a thorough review of the algorithm, including background motivation, proof of ergodicity, and illustrative examples⁸.

The hybrid method is different from the MH algorithm in how it proposes new parameter values. Instead of proposing random states based on the current value, the hybrid method uses derivatives to follow a contour of the posterior surface. By doing so, it (in theory) only proposes states that are very likely to be accepted, and as such will have less autocorrelation.

In some models, a well tuned hybrid chain will need less thinning, if any at all, and run faster. The downside of the algorithm is that it is more difficult to tune than the MH algorithm.

6.1 Algorithm

The algorithm utilizes the properties of a physical system known as Hamiltonian dynamics. Hamiltonian dynamics, while based in physics, provides some extremely useful mathematical properties for Bayesian integration via MCMC.

A Hamiltonian system consists of two parameter vectors of equal length: “position” (\mathbf{q}) and “momentum” (\mathbf{p}). How these parameters change over time is described by the Hamiltonian function, $H(\mathbf{q}, \mathbf{p})$. This system can be conceptualized as a frictionless surface about which an object moves. At some time t an object has a certain height (position) and momentum. The height of the surface is equal to the objective function of our model, and the momentum variables are introduced parameters to ensure the Hamiltonian dynamics are met. Samples from a posterior are generated by simulating the object moving about the joint surface through time, governed by H .

For use with MCMC, H is assumed to be $H(\mathbf{q}, \mathbf{p}) = U(\mathbf{q}) + K(\mathbf{p})$, where U is analogous to the potential energy and K kinetic energy. U is set equal to the ADMB objective function (i.e. the negative log of the posterior density) and K to a diagonal multivariate normal distribution. The hybrid MCMC algorithm samples from the joint posterior, H , but we are only interested in the posterior for U , so K is not saved by ADMB.

The time trajectory of the object through the joint probability space, H , is used to generate proposed parameter sets. Given the form for H above, the fundamental equations of motion are:

$$\frac{dq_i}{dt} = \frac{\partial K}{\partial p_i} \quad (2)$$

$$\frac{dp_i}{dt} = -\frac{\partial U}{\partial q_i} \quad (3)$$

ADMB uses the “leapfrog” method to discretize equations (2). The leapfrog method is more reliable than the well-known Euler method. It has two tuning parameters: the number of steps to take (`hynstep`), and the step size ε (`hyeps`). The following sequence of calculations shows a single iteration of the leapfrog

⁷Ignoring, for example, the transformation of the parameter space via the Choleski decomposition used by ADMB

⁸This chapter is available at <http://www.admb-project.org/developers/workshop/1a-jolla-2010/ham-mcmc.pdf/view>

method, for the i^{th} variable, and is repeated `hynstep` times sequentially.

$$\begin{aligned} p_i(t + \varepsilon/2) &= p_i(t) - (\varepsilon/2) \frac{\partial U}{\partial q_i} q(t) \\ q_i(t + \varepsilon) &= q_i(t) + \varepsilon \frac{p_i(t + \varepsilon/2)}{m_i} \\ p_i(t + \varepsilon) &= p_i(t + \varepsilon/2) - (\varepsilon/2) \frac{\partial U}{\partial q_i} q(t + \varepsilon) \end{aligned}$$

The leapfrog algorithm moves deterministically through the joint surface along a contour of (approximately) constant H . That is, for a given starting value of (q_0, p_0) and tuning parameters the trajectory will always end in the same place. Examples of trajectories under different tuning parameters for the leapfrog method are given in figure 5. Note that ADMB calculates the partial derivatives $\frac{\partial U}{\partial q_i}$ at each function evaluation using automatic differentiation. Thus these are available for any model and do not have to be determined analytically, which can be difficult or impossible for large, non-linear models.

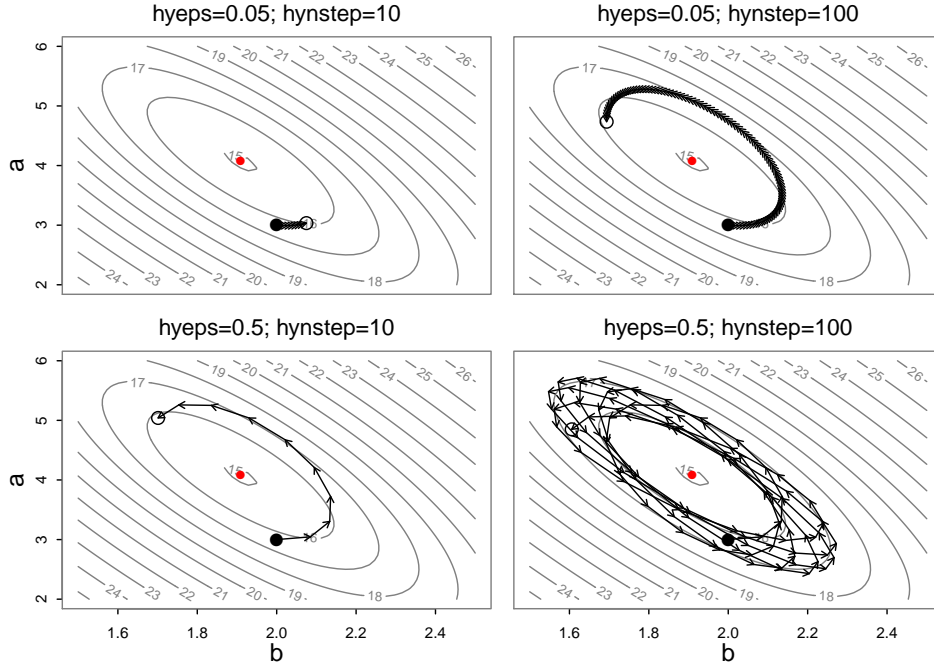


Figure 5: Leapfrog trajectories for different sets of tuning parameters. The posterior surface is shown as contours, and the posterior mode as a red point. The filled black point is the starting point, and the arrows show the trajectory of the leapfrog steps, ending at the open circle representing the proposed set of parameters.

Casting a posterior into a Hamiltonian system and discretizing it with the leapfrog method is simply a way to generate proposed sets of parameters in the larger, stochastic MCMC algorithm. A single iteration of the hybrid MCMC algorithm, as implemented in ADMB, has three steps:

1. **Propose new momentum variables.** New momentum values, p^* , are generated from a normal distribution based on the estimated covariance matrix, and independent of the current position variables.
2. **Propose new position variables.** Given the current state of the system, (q, p^*) , new position variables q^* are generated with the leapfrog algorithm using `hynstep`⁹ steps and a step size of `hyeps`.

⁹Actually a random integer is generated around this value

3. **Accept or reject the new state.** The new state is then updated with a Metropolis step (i.e. accepted or rejected) in the same way as above. Acceptance is expected to be high because the proposed values (q^*, p^*) are rarely into regions of low density for a well tuned chain.

We can see the first two steps of the algorithm by randomly drawing values for U and doing a leapfrog projection to the proposed parameters (Fig. 6).




Figure 6: Leapfrog projections for different random draws from U .

A perhaps intuitive question might be: why bother with the momentum variables at all? One issue is that without the momentum variables, and the Hamiltonian dynamics in general, we would need to account for changes in volume in the acceptance probability (step 3 above)[5]. This would require computing the determinant of the Jacobian matrix of the mapping defined by the dynamics. This Jacobian is not readily available and is often computationally intensive. Thus the need to adopt the Hamilton dynamics framework. Hamiltonian dynamics also provides other properties required for an ergodic Markov chain [5].

The hybrid MCMC thus samples from the joint posterior of the position and momentum vectors, but ADMB discards the momentum variables and returns only the position variables. The user can then do inference on those samples in the same was as the Metropolis-Hastings method, given they come from a properly converged (stationary) chain.

6.2 Arguments

Arguments for the hybrid algorithm are in general similar to those above, so we only discuss the differences. The main arguments are:

<code>-mcmc N</code>	Run N MCMC iterations
<code>-hybrid</code>	Use the hybrid method
<code>-hynstep N</code>	Mean number of steps for the leapfrog method
<code>-hyeps X</code>	The stepsize for the leapfrog method [X numeric and > 0]

Table 3: ADMB runtime arguments for the hybrid MCMC

Since the estimated covariance matrix is used in the algorithm, the `mcdiag`, `mcrb` `N`, and `mcmult` options above are also available. The `mcprobe` argument is not currently supported for the hybrid algorithm.

Note that `mcsave` `N` is **not** an argument for the hybrid algorithm and will be ignored. **Each MCMC iteration of the hybrid algorithm is saved, such that any thinning must be done manually by the user after running.**

6.3 Tuning the hybrid algorithm

A tuned hybrid MCMC algorithm often provides a more efficient (computationally) chain than the random walk behavior of the Metropolis-Hastings algorithm. However, the hybrid algorithm is often much more difficult to tune. A thorough review of tuning techniques is beyond the scope of this guide, and we refer users to [5] for further reading of more advanced tuning techniques and intuition about reasonable values¹⁰.

In most practical applications the hybrid method is tuned via trial and error. That is, values for `hynstep` and `hyeps` are tried on a short chain, and the output diagnosed. ADMB uses default values of 10 and 0.1, respectively, and these may be good starting values for most problems. The end goal is to find a set of tuning parameters that produces a well mixing chain with the fewest leapfrog steps. Figure 7 shows the autocorrelation of a parameter from the simple model across different tuning parameters. Note that the top right and bottom left panels show very similar ACF patterns, which should not be a surprise when compared to the corresponding leapfrog trajectories in 5. Further, from that figure we can see that a chain with a `hynstep` of 100 cycles the surface many times, and although it mixes well, this value could probably be lowered and the runtime reduced without affecting mixing. This is somewhat analogous to having a higher thinning rate than is necessary.

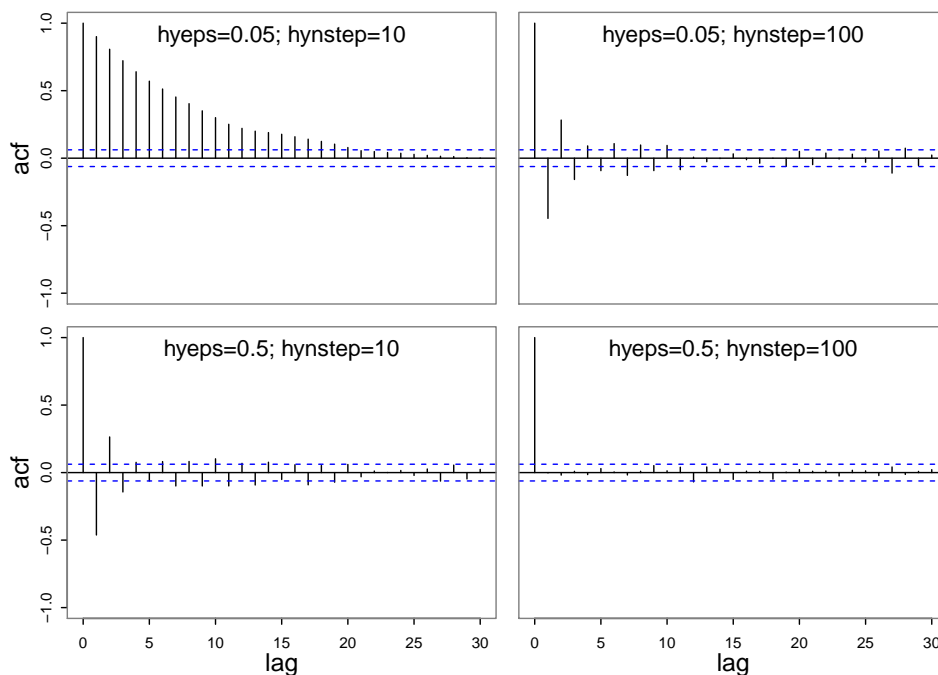


Figure 7: The autocorrelation of the parameter a from the simple model across different tuning parameters of the hybrid method.

A word of caution regarding the tuning of hybrid chains. For most values of `hyeps` the leapfrog trajectories will be “stable” in that they will cycle around a contour of H (this can be seen in the last panel in figure 5). However, for values of `hyeps` that are too large, the method becomes unstable and trajectories diverge (H is

¹⁰The ADMB algorithm works in the transformed space, via the Cholesky decomposition of the estimated covariance matrix, making intuition about tuning even more difficult

no longer constant), causing the algorithm to propose parameters with low density which are then rejected. Unfortunately for real problems, the value of `hyeps` that is “too big” can vary across the posterior space. Another issue that can arise is that certain combinations of tuning parameters can lead to near periodicity. That is, the leapfrog trajectory ends very near to where it began, after one or more steps. In pathological cases it could cycle forever and never be able to sample regions of the posterior space (i.e. not be ergodic). In practice near periodicity will make for a very slow mixing chain. ADMB mitigates this possibility by randomly drawing the number of leapfrog steps at each MCMC iteration. However, the user should still be aware of these potential issues and be vigilant in diagnosing the mixing behavior of a hybrid chain. Samples from an MCMC chain that is not in equilibrium, or has other issues, can lead to incorrect inference.

7 An example with non-linear posterior

The simple example used above had the characteristic that the parameters were correlated, but this correlation was constant across the entire posterior. Thus the same proposal distribution worked well throughout the parameter space.

Here we look at another example with two parameters, but where the correlation changes. This example uses contrived data and a discrete theta-logistic model. This is a classic example of a simple model which can lead to convergence issues for MCMC chains.

We first run it with the default settings and a thinning rate of `mcsave 100`.




Figure 8: The logistic model with 1 in 100 samples saved using the MH algorithm and estimated covariance matrix (red ellipse).

We can immediately see the autocorrelation remains high even after thinning. One option would be to increase the thinning rate. Another is to try a lower correlation. From this initial run we calculate the empirical covariance matrix and rerun the chain with this option (see 5.5.5).




Figure 9: The logistic model with 1 in 100 samples saved using the MH algorithm and empirical covariance matrix (blue ellipse). Note the improvement in acf compared to 8.

Here we see a substantial improvement in the chain performance by changing the proposal distribution of the Metropolis-Hastings algorithm. This option is far superior to increasing the thinning rate because it takes no more time to run.

We can also try the hybrid algorithm on this model, starting with the default parameters of `hyeps` and `hynstep` of 0.1 and 10, and using the estimated covariance matrix.

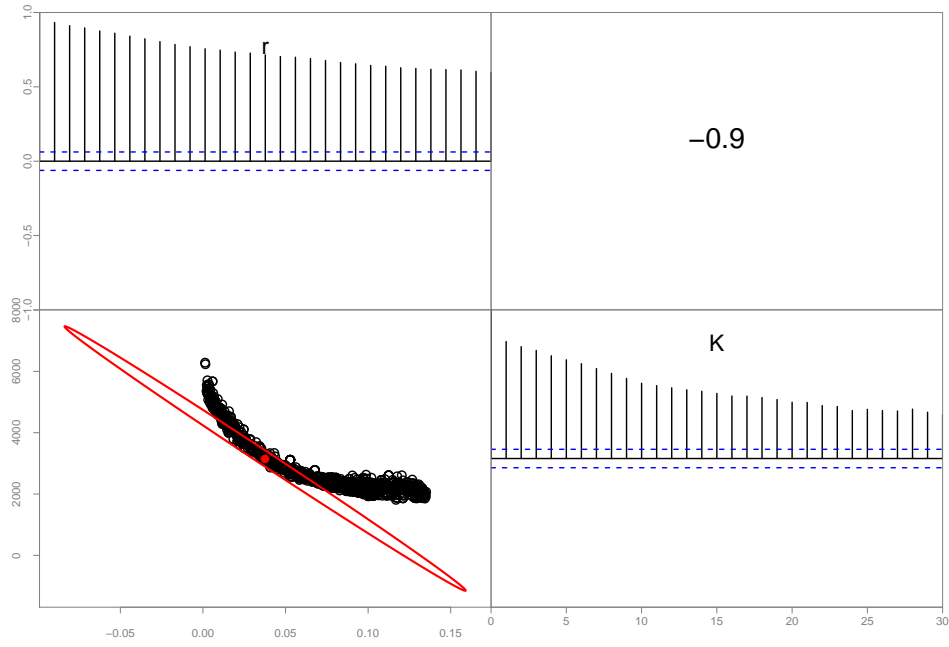


Figure 10: The logistic model using the hybrid method with default tuning parameters and the estimated covariance matrix (red ellipse).

The autocorrelation is high, suggesting we need to tune the chain. First we try the same as above, and instead use the empirical covariance matrix.

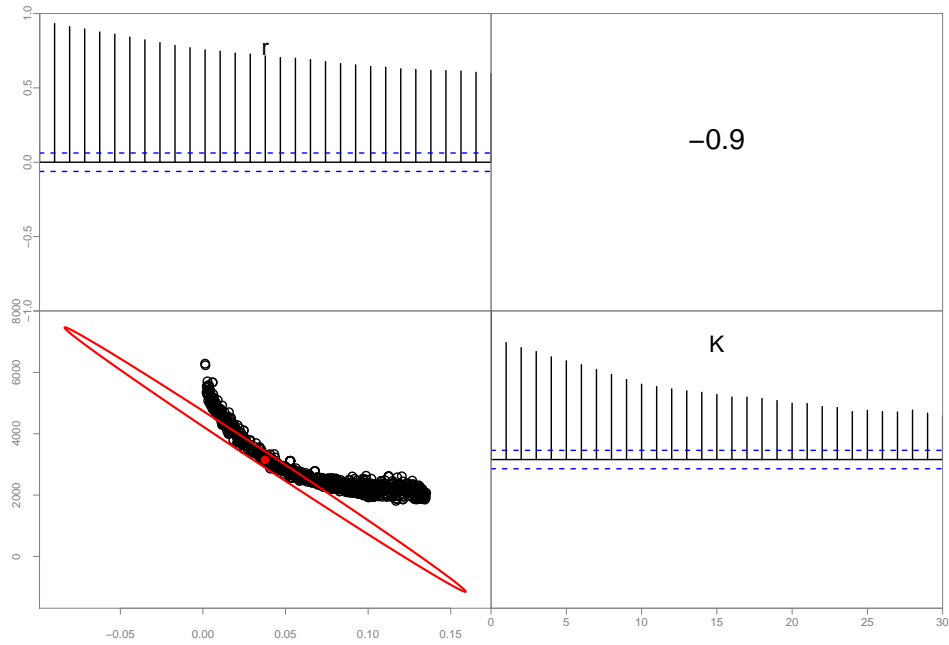


Figure 11: The logistic model using the hybrid method with default tuning parameters and the empirical covariance matrix (blue ellipse).

This chain is performing better, but is still autocorrelated. A little trial and error the combination of `hyeps 0.05` and `hynstep 100` performed significantly better, and similarly to the tuned MH chain in 9.




Figure 12: The logistic model using the hybrid method and `hyeps 0.05` and `hynstep 100` for tuning parameters and the empirical covariance matrix (blue ellipse).

Here the hybrid and MH perform similarly with the same number of function evaluations per sample (`hynstep` and `mcsave`, respectively). With a little trial and error, and visual exploration of the posterior surface it was fairly straightforward to improve on the default performance. It is likely possible to further refine the hybrid algorithm to outperform the MH, but it is not clear which of the parameters to start with. The difficulty in tuning illustrates a common issue with the hybrid method.

For many problems, the default MH algorithm will work well and require no tuning beyond a modest thinning rate. We encourage users to start with this algorithm, and explore alternatives only if computational efficiency is restrictive for the needs at hand. The hybrid method is more sophisticated, but much more difficult to tune, and we recommend it as an option only when the MH is struggling.

References

- [1] David A Fournier, Hans J Skaug, Johnnoel Ancheta, James Ianelli, Arni Magnusson, Mark N Maunder, Anders Nielsen, and John Sibert. Ad model builder: using automatic differentiation for statistical inference of highly parameterized complex nonlinear models. *Optimization Methods and Software*, 27(2):233–249, 2012.
- [2] Kimberly F. Colvin. Kruschke, j. k. (2011). doing bayesian data analysis: A tutorial with r and bugs. burlington, ma: Academic press. *Journal of Educational Measurement*, 50(4):469–471, 2013.
- [3] Martyn Plummer, Nicky Best, Kate Cowles, and Karen Vines. Coda: Convergence diagnosis and output analysis for mcmc. *R news*, 6(1):7–11, 2006.
- [4] G. O. Roberts and J. S. Rosenthal. Optimal scaling for various metropolis-hastings algorithms. *Statistical Science*, 16(4):351–367, 2001.
- [5] Steve Brooks, Andrew Gelman, Galin Jones, and Xiao-Li Meng. *Handbook of Markov Chain Monte Carlo*. CRC Press, 2011.